

2024 CSCS C++ course: modules

Created: 2024-08-22 Thu 11:25

Table of Contents

- 1. Intro: Why do we need modules? What are modules? And what are they not?
- 2. A brief overview of module structure and syntax
- 3. Modules in pika
- 4. Miscellaneous issues etc.
- 5. Summary

1. Intro: Why do we need modules? What are modules? And what are they not?

1.1. Repeated inclusion

- Headers parsed again in each translation unit
- Include guards required to avoid duplicate definitions
- `#pragma once` is not standard

1.2. Transitive includes

- Programs can compile without explicitly including everything they need
- Break when transitive includes are removed
 - https://gcc.gnu.org/gcc-13/porting_to.html#header-dep-changes

1.3. ODR: One definition rule

- Must have exactly one definition of every non-inline entity
- Can have multiple inline definitions as long as they are the same
- If you fail to follow this: undefined behaviour

1.4. Stateful includes

```
#define NDEBUG  
#include <cassert>
```

```
#define WIN32_LEAN_AND_MEAN  
#include <Windows.h>
```

```
struct foo {  
#ifndef NDEBUG  
    int debug_var;  
#endif  
};
```

```
#define private public
```

```
#include <complex.h>  
template <std::size_t I> struct;
```

1.5. Visibility

```
struct foo {  
    private:  
        void bar();  
    public:  
        void baz();  
};
```

```
__attribute__((visibility("hidden"))) void foo() {}
```


1.6. C++ modules

- C++ modules encapsulate translation units more strongly than previously
 - Names are attached to modules
- Only preprocessor definitions set during compilation of module can affect its behaviour, it can't be changed anymore when importing
- No (forced) separation of interface and implementation to hide implementation
- Import order does not matter
- No huge includes due to transitive includes
- Better compilation times
- No include guards
- No risk of missing required includes
- Modules added in C++20
- `std` and `std.compat` modules added in C++23

1.7. Modules are orthogonal to namespaces

```
export module foo;  
namespace bar {  
    void baz();  
}
```

1.8. Modules are not hierarchical

- Modules can be named `foo.bar.bar`
- The standard couldn't care less, this is simply convention
- `foo:bar` is a *module fragment* `bar` inside the module `foo`, to be covered later
 - closest thing to a submodule, but only one level

1.9. Modules have no correspondence to filenames or the filesystem in general

- Can have a module `foo` defined in a directory `bar` implemented in a file called `baz.cpp`

1.10. Modules are not simple

- Can't leave the entire header-world behind
- Have to interoperate with libraries that don't provide modules
- Need *global module fragment*
- Need includes (or header units) for macros

2. A brief overview of module structure and syntax

- Playground: <https://godbolt.org/z/qM91dTTvY>

2.1. Consuming modules is simple

```
import std;  
int main() {  
    std::println("hello");  
}
```

2.2. Writing simple modules is simple

```
// define the module interface
export module foo;
// import other modules
import foo2;
// internal definitions
void bar() { std::cout << "bar\n"; }
// exported definitions
export void bar();
export { void baz() { quack(); /* from foo2 */ } }
// internal functionality and definitions
module :private;
void bar() { bar(); }
```


2.3. Writing complicated modules is a bit more complicated

- *module unit*
- *module purview*
- *global module*
- *global module fragment*
- *module interface unit/module implementation unit*
- *primary module interface unit*
- *module partition*
- *private module fragment*

2.4. Module unit

A module unit is a translation unit that contains a module-declaration.

— **<https://eel.is/c++draft/module#unit-1>**

```
// can only put some things here
module foo;
// this is part of the module
```

```
// also a module unit
export module foo;
```

2.5. Module purview

A module unit purview is the sequence of tokens starting at the module-declaration and extending to the end of the translation unit.

— **<https://eel.is/c++draft/module#unit-5>**

```
// not module purview  
module foo;  
// module purview
```

2.6. Global module

The global module is the collection of all global-module-fragments and all translation units that are not module units. Declarations appearing in such a context are said to be in the purview of the global module.

— **<https://eel.is/c++draft/module#unit-6>**

2.7. Global module fragment

A global-module-fragment specifies the contents of the global module fragment for a module unit. The global module fragment can be used to provide declarations that are attached to the global module and usable within the module unit.

— **<https://eel.is/c++draft/module#global.frag-2>**

```
module;  
// global module fragment  
// "Prior to phase 4 of translation, only preprocessing directives  
// can appear here"  
#include <execution>  
module foo;
```

2.8. *Module interface unit/Module implementation unit*

A module interface unit is a module unit whose module-declaration starts with export-keyword; any other module unit is a module implementation unit.

— **<https://eel.is/c++draft/module#unit-2>**

```
// module interface unit  
// not a module implementation unit  
export module foo;
```

```
// module implementation unit  
// not a module interface unit  
module foo;
```

2.9. Primary module interface unit

A named module shall contain exactly one module interface unit with no module-partition, known as the primary module interface unit of the module; no diagnostic is required.

— **<https://eel.is/c++draft/module#unit-2>**

```
// primary module interface unit  
export module foo;
```

```
// can't have another module interface unit for foo  
// export module foo;
```

2.10. *Module partition*

A module partition is a module unit whose module-declaration contains a module-partition. A named module shall not contain multiple module partitions with the same module-partition. All module partitions of a module that are module interface units shall be directly or indirectly exported by the primary module interface unit ([module.import]). No diagnostic is required for a violation of these rules.

— **<https://eel.is/c++draft/module#unit-3>**

```
// module partition  
module foo:bar;
```

```
// module partition and interface unit  
export module foo:baz;
```


2.11. *Private module fragment*

A private-module-fragment shall appear only in a primary module interface unit ([module.unit]). A module unit with a private-module-fragment shall be the only module unit of its module; no diagnostic is required.

— ***<https://eel.is/c++draft/module#private.frag-1>***

```
module foo;  
module :private;  
// private module fragment
```

3. Modules in pika

- Plan: convert pika to use modules ([**https://github.com/pika-org/pika**](https://github.com/pika-org/pika))
- Probably more complicated than most other libraries, so you may get away with something simpler
- Proof-of-concept implementation: [**https://github.com/pika-org/pika/compare/main...modules**](https://github.com/pika-org/pika/compare/main...modules) (with `import std`;: [**https://github.com/pika-org/pika/compare/main...modules-import-std**](https://github.com/pika-org/pika/compare/main...modules-import-std))
 - Beware: Frankenstein branch

3.1. Prerequisites

- I used:
 - clang 18.1.7
 - cmake 3.29.3
 - ninja 1.12.1
- GCC 14 also has relatively good support, but had some issues
- Generally, the newer the better since things are being fixed rapidly

3.2. pika's existing “module” structure

- pika/
 - execution/
 - include/pika/execution/
 - algo.hpp
 - ...
 - src/
 - algo.cpp
 - ...
 - CMakeLists.txt
 - schedulers/
 - runtime/
 - ...

3.3. pika's existing “module” structure

- Each “module” compiled into an object library
- Object libraries linked into `libpika.so`
- Headers installed into single include directory from different modules
- Users only see a single library, not the individual “modules”

3.4. pika's C++ modules structure

- Each “module” becomes a C++ module: `pika.execution` etc.
- Expose a high level module called `pika.all`
 - Reexports everything
- Expose a high level module called `pika`
 - Reexports public API
- Want to keep existing headers as unchanged as possible to allow non-modules usage
- Macros are handled separately
- Mechanical translation of each “module” to a module

3.5. Step 1: defining a module

- New `module.cpp` file generated for each pika module, defines module interface

```
// Global module fragment
module;
#include <type_traits>
#include <boost/container/small_vector.hpp>
import std; // If available
// Module interface
export module pika.execution;
// Import other pika modules
import pika.config;
import pika.thread_pools;
// Export everything that we had defined in the headers
export {
#include <pika/execution/algorithms/bulk.hpp>
#include <pika/execution/algorithms/when_all.hpp>
}
```


3.6. Step 2: header files

- Only keep includes internal to the module (and macro includes) in header files

```
#pragma once // still required to avoid multiple definitions in module
unit
// Only preprocessor definitions, ok to include; could also be in the
// global module fragment or command line
#include <pika/config.hpp>
// Keep includes within the "module"
#include <pika/execution/detail/partial_algorithm.hpp>
// Don't include functionality from other pika modules; imported in
// primary module interface
// #include <pika/functional/tag_invoke.hpp>
// If included in the global module fragment, will not get included
// again; if using std module don't include
// #include <functional>
// #include <type_traits>
// #include <utility>
// Actual functionality, exported by the export block in the primary
// module interface
namespace pika::execution {
    // ...
}
```

3.7. Step 3: cpp files

- Transform cpp files to module implementation units

```
module;
// Don't include any pika headers here; declared or defined in primary
module interface
// #include <pika/execution/detail/helpers.hpp>
// #include <pika/datastructures/variant.hpp>
// #include <pika/string_util/bad_lexical_cast.hpp>
// Functionality used only in the implementation; if using std module
don't include, import std instead
#include <typeinfo>
// Module implementation unit
module pika.execution;
// Could import modules for private use here
namespace pika::execution {
// ...
}
```

3.8. Step 4: macros...

- PIKA_ASSERT, PIKA_LOG, and PIKA_VERSION etc.
- Ideal world
 - Constants become `inline constexpr` variables
 - Function-like macros use `std::source_location` and hope for inlining
- Real world
 - May need to use preprocessor to choose code paths
 - `constexpr-if` can't be used in all contexts, e.g. defining members
 - Logging/testing/assertion macros that print the expression
 - Compatibility, i.e. not having to change everything at once

3.9. Step 4: macros...

- Not all compilers support function-like macro definitions on the command line
 - CMake will remove them if set through `target_compile_definitions`, but can still pass them manually as compiler flags

```
#include <pika/assertion.hpp>
import pika;

int main() {
    PIKA_ASSERT(false);
}
```

```
// pika/assertion.hpp
// define macros only
#ifdef PIKA_DEBUG
#define PIKA_ASSERT(...) pika::handle_assertion(...);
#else
#define PIKA_ASSERT(...)
#endif
```

```
// assertion module
```

```
module pika.assertion;  
export void pika::handle_assertion();
```

3.10. Step 5: define `pika.all` module

```
module pika.all;  
  export import pika.assertion;  
  export import pika.execution;  
  export import pika.runtime;  
  // etc.
```

3.11. Step 6: define pika module

```
module pika;

import pika.assertion;
import pika.execution;
import pika.runtime;
// etc.

namespace pika {
export using ::pika::start; // Have to fully qualify names
export using ::pika::stop;
// etc.
}
```

3.12. CMake configuration

- Should be supported by CMake, meson, build2, and possibly others
 - Only tried CMake

```
cmake_minimum_required(VERSION 3.28) # non-experimental in 3.28,
import std in 3.30
project(modules CXX) # must declare language
add_library(lib)
target_compile_features(lib PUBLIC cxx_std_20) # at least C++20
target_sources(lib
    PUBLIC
    FILE_SET cxx_modules TYPE CXX_MODULES # we're building C++ modules
    FILES lib.cpp # module unit interfaces
)
target_sources(lib PRIVATE lib_impl.cpp) # module implementation units
```


3.13. Test executable

- NB. This particular example does not work on the modules branch since not all pika modules were translated

```
#include <fmt/printf.hpp>
#include <stdexec/execution.hpp>
#include <pika/assert.hpp>
import std;
import pika;

int main() {
    pika::start();
    bool result = stdexec::sync_wait(
stdexec::schedule(pika::execution::experimental::thread_pool_scheduler
{ }) |
    stdexec::then([] { fmt::println("hello"); }));
    PIKA_ASSERT(result);
    pika::finalize();
    pika::stop();
}
```

3.14. A better module implementation for pika?

- Don't export everything from every module
 - Individually export names
- “Modules” as module partitions
 - Single CMake library target; currently each “module” is a separate object library
 - Could use internals without exporting them to everyone
- Translate macros to inline constexpr variables, inline functions, or compiler flag definitions

3.15. Build times

variant	libpika/1 thread	libpika/4 threads**	test executable*
modules	35-45 s (342 targets!)	20-30 s	2.5-3.5 s
no modules	115-130 s (72 targets)	55-70 s	7-8 s
no modules (pch)	50-60 s	20-35 s	3-5 s
no modules (pch, unity)	45-50 s (31 targets)	25-30 s	3-5 s

* standalone_thread_pool_scheduler_test

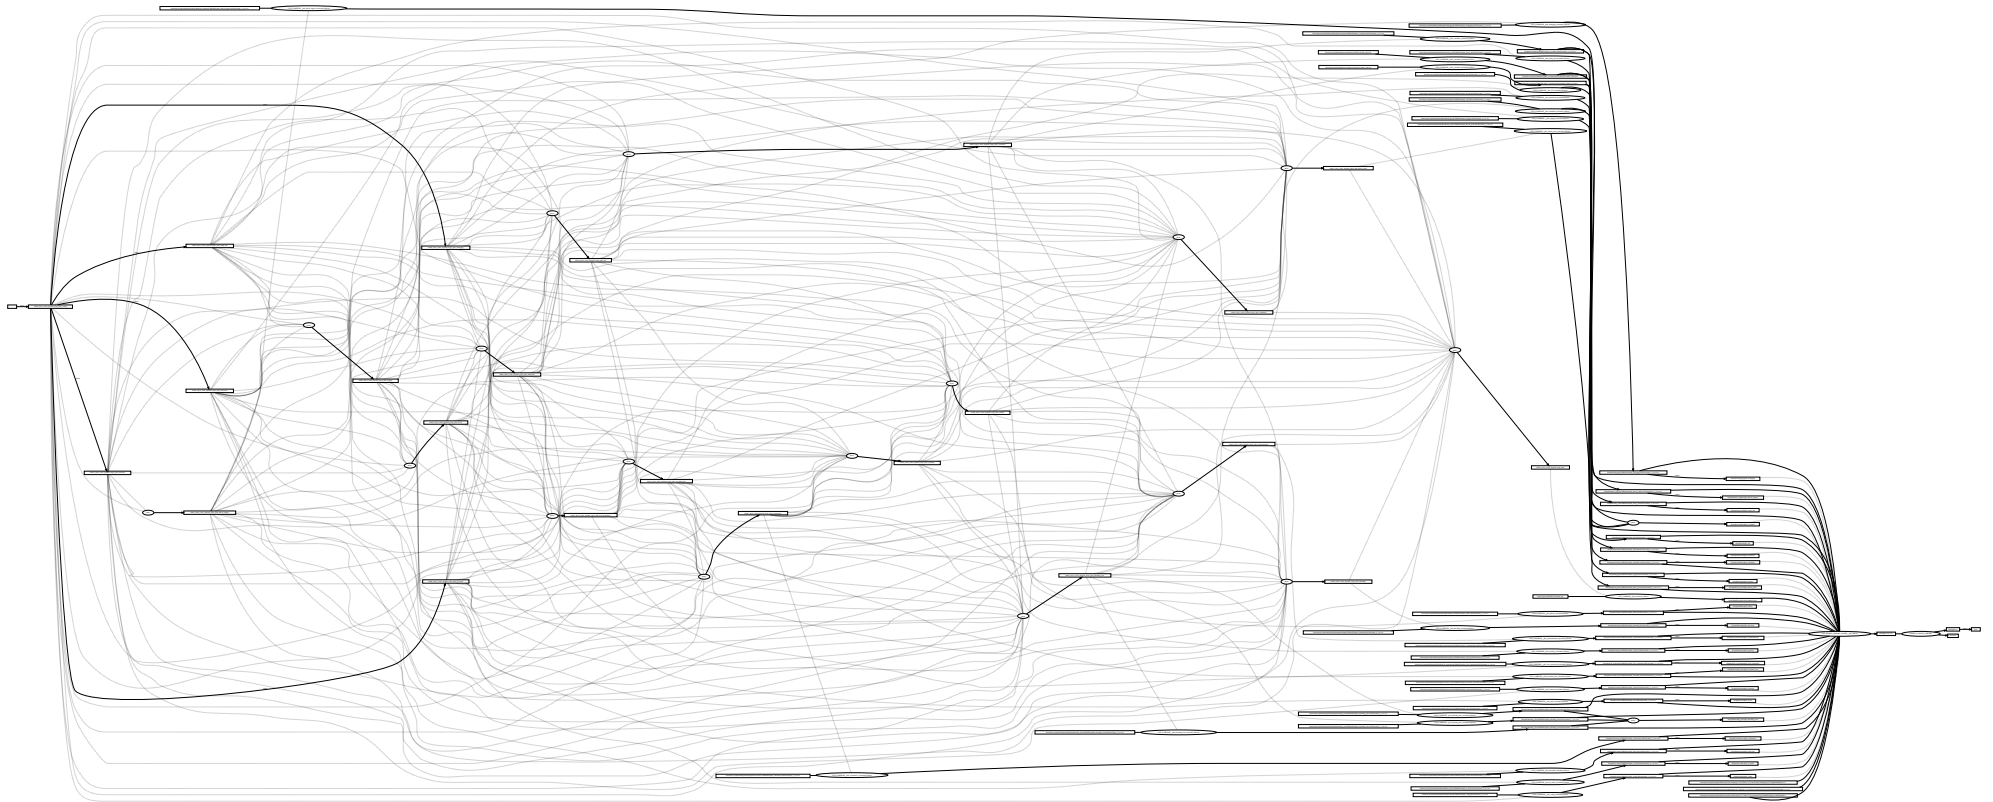
** benchmarks on noisy 4-core laptop; timing ranges from ~3 builds

3.16. Binary sizes (release mode)


variant	libpika.so	test executable
modules	1767584 bytes	414392 bytes
no modules (unity)	1831248 bytes	442560 bytes

- Possibly a small size benefit due to LTO-like behaviour with modules?
 - Too small sample size to draw general conclusions

3.17. Build graph (no modules)



3.18. Build graph (modules)



4. Miscellaneous issues etc.

4.1. CMake

- Segfault if files not in correct file set
 - Implementation (also module) files are regular source files (added via `target_sources` or `add_library`)
 - Module interfaces must be in `CXX_MODULES` file set
- CMake/ninja does not always correctly track dependencies on partial rebuilds?
 - May be fixed with newer CMake or I'm using it wrong?

4.2. GCC

- ICE including too many headers in global module fragment?
 - No particular header triggers it, but fails when including “enough” headers
- ICE including spdlog headers in global module fragment
 - **<https://bugs.launchpad.net/ubuntu/+source/gcc-11/+bug/1945364>**?
 - Disabled logging in pika’s module branch
- Does not support private module fragment

4.3. clang

- Generally good, including error messages!
- Header unit support is experimental
- Warns about includes in module purview
 - warning: `'#include <filename>'` attaches the declarations to the named module `'pika.preprocessor'`, which is not usually intended; consider moving that directive before module declaration
- Can't use macro to define module
 - macro for module mod; module not found
 - confuses dependency scanner?

```
#ifdef PIKA_HAVE_MODULE
#define PIKA_MODULE_DECLARATION(name) module name;
#else
#define PIKA_MODULE_DECLARATION(name)
#endif

PIKA_MODULE_DECLARATION(pika.execution)
```


4.4. Name mangling

- Names attached to a module `m` are suffixed with `@m`
- E.g. multiple definitions when including standard library headers in module `purview`

```
error: reference to '__and_' is ambiguous
note: candidates are: 'template<class ... _Bn> struct
std::__and_@pika.config'
note:                  'template<class ... _Bn> struct std::__and_'
```

4.5. Including headers in multiple modules

- Names are attached to a module, can't define or declare them in different modules
- No forward declarations of names from other modules
- E.g. including pika headers in different module purviews

```
pika/libs/pika/type_support/include/pika/type_support/pack.hpp:17:12:  
error: cannot declare 'struct  
pika::util::detail::pack@pika.type_support<Ts>' in a different module  
17 | struct pack  
   |           ^~~~  
In file included from pika/build/spack/libs/pika/type_support/  
module.cpp:38,  
of module pika.type_support, imported at pika/build/spack/libs/pika/  
datastructures/module.cpp:26:  
pika/libs/pika/type_support/include/pika/type_support/pack.hpp:17:12:  
note: previously declared here  
17 | struct pack  
   |           ^~~~
```

4.6. extern "C++"

- All definitions within the module purview attached to a module
- Can use extern "C++" to not attach a name to the module
- Use cases
 - Mixing modules and non-modules usage
 - CUDA/HIP translation units?
 - definition can be in a non-module translation unit

```
// fmt.cc
// If you define FMT_ATTACH_TO_GLOBAL_MODULE
// - all declarations are detached from module 'fmt'
// - the module behaves like a traditional static library, too
// - all library symbols are mangled traditionally
// - you can mix TUs with either importing or #including the {fmt}
API
#ifdef FMT_ATTACH_TO_GLOBAL_MODULE
extern "C++" {
#endif
```


4.7. static in headers

- Static functions and variables included through headers in global module fragment are not visible in module purview
 - Bug or feature?
 - Locally patched Boost to remove static
 - Generally: avoid static functions and variables in headers, prefer `inline` / `inline constexpr`

```
In file included from /pika/libs/pika/synchronization/src/detail/  
condition_variable.cpp:26:  
include/boost/intrusive/slist.hpp:580:28: error: no matching function  
for call to 'uncast'  
580 |     { return const_iterator(detail::uncast(this-  
>get_end_node()), this->priv_value_traits_ptr()); }  
    |                               ^~~~~~
```

```
template<class ConstNodePtr>  
static typename uncast_types<ConstNodePtr>::non_const_pointer  
uncast(const ConstNodePtr & ptr)
```

4.8. static in headers

- Can't use static inside export {} block

```
module foo;  
export {  
  #include <bar.hpp>  
}
```

```
// bar.hpp  
static bool bar = false;
```

```
bar.hpp: error: declaration of 'bar' with internal linkage cannot be  
exported
```

4.9. Modules require explicit import

- No more accidental transitive includes
- If compiler knows where definition is, it may tell you what to import

```
pika/libs/pika/threading/include/pika/threading/thread.hpp:40:23: error: declaration of 'function' must be imported before this declaration  
40 |         util::detail::function<void(std::exception_ptr const& e)>;  
    |                        ^
```

4.10. import std

- `import std` and `import std.compat` added in C++23
 - `std.compat`: *The named module `std.compat` exports the same declarations as the named module `std`, and additionally exports declarations in the global namespace corresponding to the declarations in namespace `std` that are provided by the C++ headers for C library facilities*
- Likely to be backported to C++20 mode by all major compiler vendors (<https://github.com/microsoft/STL/issues/3945>)
- CMake support requires 3.30, clang and libc++; did not test
 - first clang (nixos) had broken libc++ installation?
 - second clang (spack) was missing `clang-scan-deps` (used by CMake to discover module dependencies)
- Used a modified version of <https://github.com/DanielaE/std.module> to include all dependencies through a module *for testing*
 - Don't do this at home: pika effectively takes ownership of all its dependencies, making life difficult for dependees of pika
 - Equivalent to vendoring dependencies, which is generally also a bad idea

4.11. ADL with modules

- Big topic of *reachability*
- Example from <https://vector-of-bool.github.io/2019/10/07/modules-3.html>

```
// foo.hpp
template <typename T>
void do_something(T val) {
    // ...
}
```

```
module;
#include "foo.hpp"
export module acme;

template <typename T>
export void from_bulate(T item) {
    do_something(item);
}
```

```
import acme;

int main() {
    frombulate(42); // ERROR: No matching overload of `do_something`!
}
```

4.12. ADL with modules

- Big topic of *reachability*
- Example from <https://vector-of-bool.github.io/2019/10/07/modules-3.html>

```
// foo.hpp
template <typename T>
void do_something(T val) {
    // ...
}
```

```
module;
#include "foo.hpp"
export module acme;

template <typename T>
export void frombulate(T item) {
    do_something(item);
}

export void use_it() {
    frombulate(true);
}
```

```
import acme;
```

```
int main() {  
    frombulate(42); // ERROR: No matching overload of `do_something`!  
}
```


4.13. Unexplored

- Header units
 - Should make headers importable as separate translation units
 - Act similarly to precompiled headers, but built into the language
 - Macros are usable through header units
 - Not all headers can be imported
- Built Module Interface (BMI)
 - Installing and consuming libraries with modules
- clang/GCC compatibility
- CUDA/HIP compatibility
- etc.

5. Summary

- Modules are
 - Useful
 - Immature
 - Complex, when used in the real world
- **Should you use modules? Yes, with reservations**
 - Executables where noone depends on you
 - Then again, someone has to start using modules in libraries
 - No CUDA/HIP? clang to the rescue?
 - Otherwise, wait for better times

5.1. Resources

- <https://eel.is/c++draft/module>
- <https://en.cppreference.com/w/cpp/language/modules>
- <https://clang.llvm.org/docs/StandardCPlusPlusModules.html>
- <https://gcc.gnu.org/wiki/cxx-modules>
- <https://www.youtube.com/watch?v=iMNML689qIU>
- <https://vector-of-bool.github.io/2019/03/10/modules-1.html>
- <https://vector-of-bool.github.io/2019/03/31/modules-2.html>
- <https://vector-of-bool.github.io/2019/10/07/modules-3.html>
- <https://learn.microsoft.com/en-us/cpp/cpp/modules-cpp?view=msvc-170>
- <https://devblogs.microsoft.com/cppblog/moving-a-project-to-cpp-named-modules/>
- <https://cmake.org/cmake/help/latest/manual/cmake-cxxmodules.7.html>
- <https://www.kitware.com/import-cmake-c20-modules/>
- <https://www.kitware.com/import-std-in-cmake-3-30/>

- **<https://wg21.link/p1441>**
- **<https://wg21.link/p1689>**
- **<https://wg21.link/p1788>**

Speaker notes