




CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



std::BLAS

C++ course

Raffaele Solcà, CSCS

September 02, 2024

28.9 Basic linear algebra algorithms

<https://eel.is/c++draft/linalg>

- What is covered?
 - BLAS Level 1
 - BLAS Level 2
 - BLAS Level 3

- What is not covered?
 - LAPACK
 - PBLAS
 - ScaLAPACK

Example: GEMM

- CBLAS API:

- `cblas_gemm` (layout, transa, transb, m, n, k, alpha, a, lda, b, ldb, beta, c, ldc);

Provided by mdspan

$$C \leftarrow \beta C + \alpha \text{Op}(A) \text{Op}(B)$$

Also provided by
extensions of mdspan

- C++ API

- `std::linalg::matrix_product(A, B, C);`

$$C \leftarrow AB$$

- `std::linalg::matrix_product(A, B, E, C);`

$$C \leftarrow E + AB$$

mdspan: new features

- Some addition to `std::mdspan`
 - `std::linalg::conjugated(A);`
 - `std::linalg::transposed(A);`
 - `std::linalg::conjugate_transposed(A);`
 - `std::linalg::scaled(alpha, A);`
 - Note: difference of `std::linalg::scaled` v.s. `std::linalg::scale`
 - They return a read-only `mdspan` with a different type, layout and accessor. No operations performed.
- C++ API covers all the `cblas` GEMM cases, e.g.:
 - `matrix_product(scaled(alpha, transposed(A)), B, scaled(beta, C), C);`
- and even more (mixed precision, ...).

Can we get rid of all extra parameters?

- Consider Hermitian/symmetric matrix product

Provided by mdspan

- `cblas_hemm(layout, side, uplo, m, n, alpha, a, lda, b, ldb, beta, c, ldc);`

Different overload or Template specialization

- C++ API:

- Left:
`hermitian_matrix_product(A, t, B, C);`
`hermitian_matrix_product(A, t, B, E, C);`
 - Right:
`hermitian_matrix_product(A, B, t, C);`
`hermitian_matrix_product(A, B, t, E, C);`

There is still “diag”

- Triangular solver Need different name
Provided by mdspan
 - `cblas_trsm(layout, side, uplo, trans, diag, m, n, alpha, a, lda, b, ldb);`

Different overload or Template specialization

- C++ API:
 - Left:
`triangular_matrix_matrix_left_solve(A, t, d, B)`
`triangular_matrix_matrix_left_solve(A, t, d, B, divide)`
 - Right:
`triangular_matrix_matrix_right_solve(A, t, d, B)`
`triangular_matrix_matrix_right_solve(A, t, d, B, divide)`

Warning:
upper/lower_triangle_t refers to the
input matrix after any **transpose()**
operation has been applied.

BLAS: ‘L’, ‘T’
CBLAS: **CblasLower**, **CblasTrans**
C++: **transpose(A)**, **upper_triangle**

Summarizing

- matrix (pointer, size, leading dimension, transpose, conjugate_transpose): mdspar
- scaling factors: linalg::scaled
- side: overload or different function name
- uplo and diag:
 - ```
struct upper_triangle_t;
inline constexpr upper_triangle_t upper_triangle;
struct lower_triangle_t;
inline constexpr lower_triangle_t lower_triangle;

struct implicit_unit_diagonal_t;
inline constexpr implicit_unit_diagonal_t implicit_unit_diagonal;
struct explicit_diagonal_t;
inline constexpr explicit_diagonal_t explicit_diagonal;
```

# BLAS v.s. std::linalg

- BLAS API has only runtime parameters
- C++ API:
  - matrix sizes and leading dimensions can be decided at compile time or at runtime,
  - scaling factors at runtime,
  - BLAS side, uplo, trans and diag only at compile time.

Beneficial for small sizes.

Want possibility to decide at runtime?  
=> requires multiple instantiations of  
the code, therefore larger libraries /  
executables.

Massive template usage might  
lead to longer compilation time.



# First example

- A simple implementation of Cholesky.

<https://github.com/eth-cscs/cpp-course-2024/blob/main/stdBLAS/examples/cholesky.cpp#L57-L109>

# Quick look to reference GEMM

[https://github.com/kokkos/stdBLAS/blob/main/include/experimental/\\_p1673\\_bits/blas3\\_matrix\\_product.hpp#L655-L732](https://github.com/kokkos/stdBLAS/blob/main/include/experimental/_p1673_bits/blas3_matrix_product.hpp#L655-L732)

# Performance

- Don't expect any performance from the reference implementation.
  - An attempt to include a BLAS call to gemm has been commented out.
    - Hopefully calls to BLAS will become available.
    - However not all the calls can be mapped to BLAS. Be aware of performance penalties.
- Reference implementations are bad. Do not even include basic optimizations such as loop reordering.
- Reference implementations are not pre-compiled. Compiling your code with **no optimizations** (e.g. for debugging) means really bad performance.

# Some experiments with GEMM

<https://github.com/eth-cscs/cpp-course-2024/blob/main/stdBLAS/examples/gemm.cpp#L528-L594>

- Reference GEMM

<https://github.com/eth-cscs/cpp-course-2024/blob/main/stdBLAS/examples/gemm.cpp#L59-L62>

- Loop reordering 3 cases. Best inner loop:

- Column major:
  - NN: i, NT: i, TN: k, TT: no
- Row major:
  - NN: j, NT: k, TN: j, TT: no

<https://github.com/eth-cscs/cpp-course-2024/blob/main/stdBLAS/examples/gemm.cpp#L64-L190>

- Calling BLAS

<https://github.com/eth-cscs/cpp-course-2024/blob/main/stdBLAS/examples/gemm.cpp#L330-L376>

# Summary

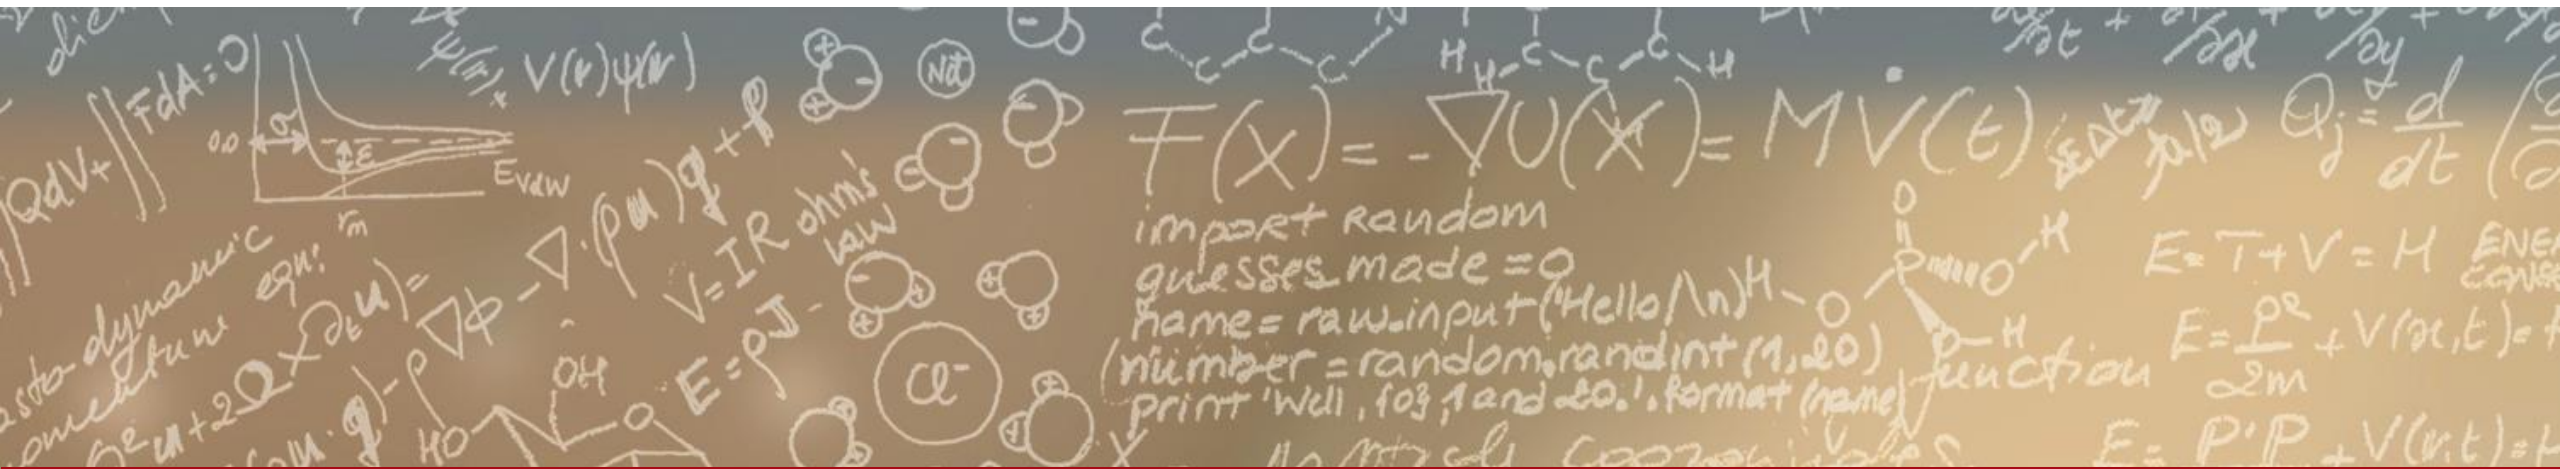
- Performance not on par with expectations.
- Compilation without optimizations: 100x – 250x slower.
- Adding basic optimizations such as loop re-ordering is not straightforward.
  - Different layout (or transpose) might need a different variant.
- The standard is not finalized yet.
  - Fixed BLAS3 rank-k reference implementation (was completely wrong).
  - Found errors in the accepted standard wording.
  - Some proposals to refine the interface are still open:
    - <https://isocpp.org/files/papers/P3371R0.html> BLAS3 rank-k and rank-2k API (breaking) change
    - <https://isocpp.org/files/papers/P3050R1.html> Optimize conjugated for non-complex types
    - <https://isocpp.org/files/papers/P3222R0.html> Introduce layout-[left, right]-padded



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



**Thank you for your attention.**